Does the books' order make the weekly review explain better?

One sentence added to the weekly and monthly routines' prompts, replayed on nine past periods

Chess Coach research lab

Abstract

The app's weekly and monthly routines have a model write a review from a brief the app's services compute. Two chess books explain a position in one order: the threat, then why the tempting move fails, then what the better move keeps. One sentence asking for that order was added to the live prompts, and both prompts rewrote nine past periods (seven weeks and two months) offline with the app's code at commit `fa07ea1ed8`: 25 reviews, for $225 of model calls. The sentence changed the writing: 79% of the candidate's board exhibits gave all three parts in order, against 32% under today's prompt. A blind judge, reading each pair in both orders, preferred the candidate's explanations in seven of twelve pairs and today's in one, but its interval reached zero on the train periods and on the held-out ones, so the rule fixed before the runs keeps today's prompts. The same judge sided with whichever review it read second in 63% of the verdicts that picked a side, and two runs of today's prompt on one period split on explanation in two of four periods. On the held-out periods the candidate's reviews made more false claims, and most of its extra false claims quote a line of moves. Two errors do not depend on the prompt. Of the 76 lines the reviews tag *Played*, 18 are not what the game played, and the app's checker, which replays each line only for legality, lets them through. And the reviews call a rule about the king in a won endgame "the pawn break", because the rule ledger files it under that concept.

1 The question

Levy Rozman's *How to Win at Chess* (2023) and *How to Win at Chess: Next Level* (2026) answered twenty questions about the app's coaching. Asked how a coach should explain a better move, they gave an order their worked examples share: first the threat on the board, then why the move played (or the tempting one) fails, with the reply that refutes it, then what the better move keeps. They do not prescribe one sentence order or a length, and their examples start in different places, so the candidate tests the order they share most, as one instruction added to step 4 of the weekly and monthly prompts, right after the sentence that says to write the annotations by filling the brief's slots:

> Each board that shows a better move explains it in its prose in this order: first the threat on the board, the piece and the square it aims at; then why the move played, or the tempting one, fails, with the reply that refutes it; then what the better move keeps, the square, piece or line it holds.

Nothing else in the prompts changes. Measured beside the comparison, for both prompts: which false claims the weekly and monthly reviews make.

2 The cases and the writer

The cases are the nine periods with a review on the app's main branch from its command-line era: weeks 35 to 41 and the months of August and September 2026. Each is a checkout of the app at `fa07ea1ed8`, with the coach data as main held it just before that period's review landed, and the period's brief prepared again by the app's own `coach weekly prepare` or `coach monthly prepare`, offline. A week holds 19 to 59 games and the months 307 and 172. The split was fixed before any run, interleaved by date: train is week 35, week 37, week 39, week 41 and August; held out is week 36, week 38, week 40 and September.

The writer runs as the routines do: Opus 5.5, headless, with nothing inherited from the session that launched it, and a short preamble that names the period, says the brief is ready and stops the writer once the strict check passes (no video, publish or commit). Each prompt wrote every case once, and again where the budget allowed a second pair (week 35, August and week 36); today's prompt also wrote week 37 twice. Run k of the candidate is judged against run k of today's prompt, and today's two runs of a period against each other give the judge's noise floor (week 35, week 37, August and week 36).

3 The graders

A script with no model reads each run: whether the final strict check passes, every check the writer ran, commands it should not have run, cost and minutes. Sonnet 5.5 audits each review alone. It reads the evidence (the brief without its bulk, a line per game, the clock's turning points, the earlier reviews the period's review audits, and for every board exhibit the position as the archive replays it, what was played there and next, the engine's best move and reply, and what each piece attacks and guards before and after each move the exhibit names) and lists the claims the evidence contradicts (*false*), does not carry (*unsupported*) or carries but slants (*misleading*). For each exhibit it says whether the prose names the threat, the refuting reply and what the better move keeps, in that order, and lists its false claims about the board. Its prompts were settled on one run of today's prompt before any candidate run was graded. Opus 5.5, tried as the auditor on that run, did not list the same false claims, so a review's count is a floor on what it gets wrong, compared like for like.

Sonnet 5.5 also judges each pair blind, in both orders, on five criteria: truth, explanation, the prompt's duties, plain words, and overall. A side wins a criterion only when both orders agree. Two controls check the judge on every case: a review of today's prompt against itself, and against a decoy whose exhibits restate what they already say at the candidate's length, with no square, move, piece or number the original does not name.

4 The rule

A pair scores $+1$ when the candidate wins a criterion, -1 when it loses and 0 for a tie, and a case scores the mean over its pairs; intervals are 95% percentile bootstrap intervals over cases. The candidate replaces the live prompts only when every condition holds (Table 1): its explanation interval lies above zero on the train cases and on the held-out ones, the overall interval does not lie below zero, neither the false claims nor the false board claims per review rise by an interval above zero, the strict check passes as often, and the controls are clean. With four or five cases a split, the first two need the candidate to win nearly every case; that is the bar. The two that fail are

explanation on train above 0 and explanation held out above 0, so today's prompts stay.

Condition	Estimate (95% interval)	Holds
Explanation on train above 0	+0.40 (−0.40 to +1.00)	**no**
Explanation held out above 0	+0.38 (0.00 to +0.75)	**no**
Overall not below 0	+0.22 (−0.17 to +0.56)	yes
False claims not up	+0.78 (−0.56 to +2.22)	yes
False board claims not up	+1.22 (−1.11 to +3.56)	yes
Check passes as often	–	yes
Controls clean	–	yes

Table 1: The rule, fixed before the candidate's first run. Explanation and overall are the judge's scores per case; the claim rows are the audit's counts per review, candidate minus today's prompt, over all nine cases.

5 What the sentence changed

The writer did what the sentence asks. Pooled over the exhibits, 79% of the candidate's 42 gave all three parts in order, against 32% of today's 44. The threat was named in 88% (today 68%), the refuting reply in 79% (70%), and what the better move keeps in 83% (73%). Every review passed the strict check in the end, 13 of 13 and 12 of 12, after 5.5 and 6.3 check rounds on average. A weekly review cost $6.62 under today's prompt and $7.55 under the candidate, and a monthly one $5.78 and $5.49; a review took about 21 minutes either way. The monthly reviews show few boards: today's three carried 0, 0 and 1 and the candidate's 1, 1 and 2, so the explanation criterion rests on the weeks (Table 2).

			Exhibits in order		Judge, per pair	
Period	Split	Games	Today	Candidate	Explanation	Overall
Week 35	train	55	3/9	7/9	+ +	+ +
Week 37	train	19	4/9	3/4	−	−
Week 39	train	59	1/4	3/4	+	+
Week 41	train	34	3/4	3/4	0	0
August	train	307	0/0	2/2	+ +	0 +
Week 36	held out	32	1/8	6/8	0 +	0 +
Week 38	held out	49	0/4	3/4	+	0
Week 40	held out	33	2/5	4/5	0	0
September	held out	172	0/1	2/2	0	0

Table 2: Each period: its games, the exhibits that give all three parts in order (over all of a prompt's runs of the period), and the judge's verdict on each pair, from the candidate's side (+ it wins, − it loses, 0 a tie).

6 What the judge saw

On explanation the judge preferred the candidate in seven of the twelve pairs, today's prompt in one and neither in four. That is +0.40 (−0.40 to +1.00) over the train cases and +0.38 (0.00 to +0.75)

over the held-out ones. Today's prompt took week 37 and week 40 and September tied, so neither interval clears zero. Overall the candidate won five pairs, lost one and tied six (+0.22 (−0.17 to +0.56) over all nine cases), and the intervals of truth (+0.11 (−0.22 to +0.44)), duties (0.00 (−0.33 to +0.33)) and plain words (0.00 (0.00 to 0.00)) hold zero.

Two things say the judge is a weak instrument for this question. Today's prompt against itself, two runs of one period, gave a decisive explanation verdict in two of four periods (0.00 (−0.75 to +0.75)), against eight of twelve for the candidate's pairs. And across all pairs and criteria, 63% of the 130 single-order verdicts that picked a side picked the review it read second; requiring both orders to agree turns most of those into ties. The controls are clean: the same review on both sides tied on every criterion in nine of nine cases, and the decoy, whose exhibits ran 1.21 times as long as the original's, won explanation in zero of nine cases and lost overall in four. The judge does not reward length alone, but with nine cases it cannot tell this sentence from the writer's run-to-run variation.

7 What the reviews get wrong

Per review, today's prompt made 2.1 false claims, 9.9 unsupported and 5.5 misleading ones, and 2.5 false board claims; the candidate made 2.9, 8.5, 4.8 and 3.6. Candidate minus today's, the false claims are −0.8 (−1.6 to +0.2) on the train cases and +2.8 (+1.5 to +4.0) on the held-out ones, where every case went up; over all nine they are +0.8 (−0.6 to +2.2), inside rule 4.

A second blind Sonnet call sorted every listed claim by what it is about, and marked an entry whose own evidence says the claim holds: 13% of the claims listed in words, and 33% of the false board claims, which Table 3 leaves out. The false claims are mostly about who stood better, the best move, a skill, metric or focus reading and results, records, ratings (61% of them); the unsupported ones spread over the brief's readings, trends and what earlier reviews said, where the evidence the auditor read may lack what the writer read.

	Today's prompt		Candidate	
What the claim is about	False	Unsupported or misleading	False	Unsupported or misleading
Who stood better, the best move	0.5	1.1	0.9	1.3
A skill, metric or focus reading	0.3	3.1	0.6	1.9
Habits and trends	0.2	2.3	0.5	2.1
Results, records, ratings	0.4	1.5	0.3	1.2
What an earlier review said	0.2	2.1	0.3	1.6
What a piece does on the board	0.2	1.0	0.3	0.7
Counts and tallies	0.2	1.8	0.0	1.0
Other	0.0	0.5	0.1	0.4
The clock	0.1	0.6	0.0	0.4
Every kind	2.0	13.9	2.9	10.5

Table 3: Claims per review the audit lists, by what they are about, without those whose evidence says they hold.

Lines of moves. False claims that quote a line of two numbered moves or more run 0.15 a review under today's prompt and 1.00 under the candidate: 0.8 of the candidate's 0.9 extra false claims a review. The sentence asks for the reply that refutes the move; the writer answered with lines, and wrote some of them wrong.

Lines tagged Played. A board exhibit carries lines with a tag, most often *Played* and *Better*. A script with no model replayed every Played line from the exhibit's position and compared it with the game in the archive. Of the 76 Played lines, 18 leave the game (18% under today's prompt, 31% under the candidate), in 13 of the 25 reviews, and so do two of the 27 in the reviews that landed. The usual shape is the move played followed by the engine's refutation, or a later move the game never had. The strict check replays each line for legality from the board's position but does not compare a Played line with the game, so none of these was refused.

The king walk called the pawn break. The weekly review reports the rule ledger's budget, which caps the live rules per concept. In weeks 36 to 41, the only live rule under the concept "Pawn structure and the break" is about the king in a won endgame. Of the 15 weekly reviews of those weeks, seven (five of eight under today's prompt, two of seven under the candidate) count it as "the pawn break", and the auditor marks five of them false; the rest say "pawn structure". The writer follows the prompt, which names a concept by its name. The claim is false because the ledger files the rule under a concept it is not about.

8 What the app takes

The live weekly and monthly prompts stay. Two fixes do not depend on the prompt: a strict check that a line tagged Played follows the game from the board's position (the checker already resolves the board's game and replays the line), and the king-walk rule filed under converting, or the per-concept count naming each rule by what it says. If the order is tried again, it would want the engine's refutation in the brief for each exhibit, so the line the writer shows is one the engine played out (an inference from the lines above, not tested), and a judge that compares the same exhibit under both prompts, which needs the exhibits fixed before the writer runs.

9 Limits

Nine cases, and at most two runs of a prompt on one; the judge's verdicts on the same prompt vary nearly as much as between prompts, and it leans to the review it reads second. The audit's counts are a floor, and a share of what it lists holds by its own evidence. One writer model wrote every review, offline: the level falls back to the profile synced at the time, and the brief is the one the services compute at `fa07ea1ed8`, not always the one the period's routine read. The writer eval spent $168 and the graders $58 at list price; on the owner's subscription the seven-day usage window moved from 12% to 25% while it ran.